Part II. Supporting Report

1. Problem and Approach

On April 24, 2026, Governor Janet Mills vetoed L.D. 307, the legislation that would have made Maine the first state to enact a temporary moratorium on data centers exceeding 20 megawatts of contracted load. The bill passed the Legislature with bipartisan support (House 79–62, Senate 21–13) but fell short of the two-thirds margin needed to override. The veto coincided with Mills's signature on L.D. 713, which removes data centers from Maine's business equipment tax exemption and Dirigo incentive programs, and with her announcement of an executive order convening a study commission similar to the one L.D. 307 itself would have established.

The political record indicates that the impasse wasn't the underlying problem. The Governor's own AI Task Force concluded in November 2025 that residents harbored substantive concerns about data center impacts on rates and the environment, and L.D. 307's sponsor (Representative Melanie Sachs) had collaborated with the Mills administration on the bill's initial drafting. The disagreement turned on a single fact pattern: a proposed redevelopment of the Androscoggin Mill in Jay, where Mills sought a project-specific carveout that the House rejected 115–29. The Governor then vetoed the underlying bill.

This proposal accepts the impasse as a constraint and routes around it. A moratorium that fails to clear the override threshold has no governance effect. A regime that the Governor can sign, and that creates durable visibility into how compute is being deployed in the state, is a more enforceable instrument than a moratorium that the Legislature cannot enact.

The proposal also takes seriously a technical observation: the proximate driver of the data center buildout is not data centers in the abstract but specifically demand for compute associated with artificial intelligence training and inference. BloombergNEF revised its projection of U.S. data center demand upward by thirty-six percent (to 106 gigawatts by 2035) within seven months in late 2025, on the basis of changes in the AI compute trajectory. Governance instruments calibrated to AI workloads specifically can therefore be both narrower in scope and more directly responsive to the underlying concern than instruments calibrated to data centers generically.

Three considerations recommend a disclosure regime over a categorical ban. The first is information asymmetry. The proximate failure that motivated L.D. 307 was that operators (per Sachs's testimony) had refused to engage with stakeholders. Under current law, an operator seeking a 100-megawatt interconnection is not obligated to disclose whether the workload is AI training, AI inference, or cryptocurrency mining, although these workloads have substantially different rate, water, employment, and policy profiles. A disclosure regime addresses this directly. The second is optionality. A moratorium is binary; a disclosure regime produces a continuous record on which the Legislature can later escalate, on the basis of evidence rather than projection. The third is the ISO-NE structure. A unilateral Maine moratorium displaces development to neighboring states without altering aggregate regional load, while disclosure requirements can be replicated across states, and Maine is well positioned to be the first mover.

None of this implies that a moratorium is wrong in principle, only that the disclosure regime is more enforceable under current Maine political conditions while preserving the option to escalate.

2. Mechanism

The Act has three operational components, each doing work the others cannot.

The Workload Composition Report (part 4) is the central informational instrument. It requires the operator to disclose, quarterly, the share of facility accelerator-hours allocated to AI training, AI inference, cryptocurrency mining, and other workloads, plus tenant-level use cases for any tenant exceeding ten percent of load and notification of any frontier training run on premises. Standard ASHRAE 90.4 metrics and any power purchase or curtailment arrangements are reported alongside. Reports are published within thirty days, and aggregate workload figures, total consumption, and frontier-run notifications are nonconfidential by statute.

The tiered review (part 5) imposes proportionate procedural burdens at the points where rate and grid impacts begin to matter. At 20 megawatts (the threshold inherited from L.D. 307 and from the New York PSC's analogous cryptocurrency proceeding), expedited review and a Grid Impact Assessment are required. At 100 megawatts, full proceedings with public hearing are required. Approval is conditioned on a finding that residential and small commercial rates will not rise as a result of the load addition.

The Grid Stewardship Agreement (part 6) is the financial and operational counterpart to disclosure. Fifty cents per megawatt-hour funds grid hardening, demand response, and existing low-income ratepayer assistance. Mandatory enrollment in a curtailment program of up to fifteen percent during declared emergencies ensures the facility contributes to grid stability rather than purchasing exemption from it. The regional siting requirement on matched generation prevents Maine from absorbing the load while the offsetting capacity is procured outside the ISO-NE control area.

3. Technical Grounding

Three technical assumptions underwrite the proposal: that AI compute can be measured at the facility level, that the training-versus-inference distinction is meaningful, and that the 10^{25}-FLOP threshold is a defensible inflection point.

On measurement: data center operators routinely instrument their facilities for billing, tenant chargeback, and energy management compliance under ASHRAE 90.4, ENERGY STAR for Data Centers, and ISO 50001. The marginal cost of adding a workload-type column to existing accounting is small, and the operator's own internal estimates need only be reported with a stated margin of estimation. The Act does not require new measurement infrastructure.

On training versus inference: training workloads tend to run as concentrated, high-utilization campaigns (often weeks of sustained near-peak load on dedicated clusters), with diurnal patterns that align poorly with renewable generation. Inference workloads tend to be more distributed, more interruptible, and more responsive to demand-following arrangements. The capability-development externalities that drive policy debate attach to training rather than inference, and the distinction therefore yields information that the regulator and Council cannot otherwise obtain.

On the 10^{25}-FLOP threshold: the cutoff is consistent with 15 C.F.R. part 744.23 (U.S. export controls), Article 51 of the EU AI Act (which uses 10^{25} FLOP as its presumption of systemic risk), and California Senate Bill 53. Section 8(b) directs the Commission to adopt a calculation convention consistent with these jurisdictions, preserving interoperability and reducing compliance cost for operators already reporting elsewhere. The threshold is acknowledged to be imperfect; its function in this Act is administrative (it triggers a notification, not a substantive determination), and the Council's annual findings provide the mechanism by which it can be revised.

The principal numerical thresholds are inherited rather than chosen de novo. The 20-megawatt covered-facility threshold tracks L.D. 307 and the New York PSC's Case 22-E-0064; it corresponds approximately to the load of a 28,000-accelerator AI cluster at current GPU power draws. The 100-megawatt threshold for full proceedings corresponds to roughly one-third of one percent of ISO-NE's installed 30-gigawatt generation base, the order of magnitude at which interconnection studies routinely identify meaningful transmission constraints. The fifty-cents-per-MWh contribution generates approximately $329,000 per year at a 100-megawatt facility operating at 75 percent utilization (about 0.5 percent of energy spend at current ISO-NE wholesale prices), calibrated to be administratively visible without deterring siting.

4. Failure Modes

Three failure modes warrant explicit treatment.

Misclassification. An operator could underreport the AI training share by reclassifying training workloads as "other," particularly where the relevant tenant is internal. The Act addresses this through three layers: tenant-level use case disclosure for any tenant exceeding ten percent of load (part 4(a)(3)), the biennial independent audit on a sample basis (part 8(a)), and a civil penalty of $10,000 per violation per day for materially false reporting (part 8(c)). The audit operates on the IRS / SEC model, where the disciplining effect derives from the credible threat of selection rather than from universal review.

Structural avoidance. An operator could attempt to evade the regime by structuring a single development as multiple sub-20-megawatt facilities at one site, by siting just across the New Hampshire border, or by contracting compute capacity from a non-Maine facility while marketing the workload as "Maine-based." The first concern is addressed by the definition of covered facility (part 3(a)), which captures any colocated facility at a single point of common coupling. The second is the structural disadvantage of any single-state regime in a regional grid; it is addressed partially through part 7(b)(iv), which directs the Council to coordinate with the New England States Committee on Electricity. The third belongs to federal supply-chain regulation and is largely outside the Act's scope.

Enforcement underinvestment. The Public Utilities Commission, funded principally through utility assessments, may lack technical capacity to evaluate AI-specific reports. The biennial independent auditor (part 8(a)) is designed to address this gap, and the Resilience Fund's penalty deposits (part

8(c)) are designed to provide a self-financing source for its administration. A small additional appropriation for the first three years would be helpful, although the structure does not depend on it.

5. Feasibility and Institutional Realism

The proposal is institutionally and legislatively conservative. It vests authority in the Public Utilities Commission, which already exercises plenary jurisdiction over interconnection within Maine, rather than in a new body. It modifies the Coordination Council that L.D. 307 contemplated and that Mills's executive order will in any event create. It uses the existing Office of the Public Advocate, the Department of Environmental Protection, and the Climate Council as Council seats. The reporting and audit instruments are built on top of standards (ASHRAE 90.4, ENERGY STAR, ISO 50001) that operators already comply with for unrelated purposes.

The proposal is structured as a substitute for L.D. 307 rather than an independent measure, preserving the rate protection, environmental review, and host-community consultation goals that produced L.D. 307's bipartisan passage. The Governor's stated objection concerned the project in Jay; part 5 of the substitute does not categorically prohibit such projects but conditions their approval on rate-impact findings and on disclosure and stewardship compliance. The substitute is therefore plausibly signable under the political conditions that produced the April 24 veto. The two methodology rulemakings required under part 8(b) are aggressive but achievable in the 180-day window; the New York PSC completed an analogous rulemaking for cryptocurrency loads in a comparable window in 2022.

6. Limitations

Three limitations warrant acknowledgement. First, the Act addresses governance gaps at the infrastructure layer rather than at the model development, training data, or downstream deployment layers, and is best understood as one of several layered interventions. Second, the Act's targeting depends on AI workloads remaining a material share of data center capacity; the Council's annual review provides the mechanism by which targeting would be revisited. Third, no enforcement mechanism can fully secure operator candor against a determined and well-counseled adversary. This is a generic limitation of any disclosure regime, and it is the principal reason that the substitute does not foreclose the option of a subsequent moratorium. The disclosure regime is the first move in a sequence, not the last.

7. Conclusion

The April 24 veto record establishes the political conditions under which AI compute governance must be enacted in Maine: an underlying public concern that has cleared a bipartisan supermajority of legislative opinion, an executive willing to sign incremental measures but not categorical prohibitions, and a substantive informational gap that no party has defended. The substitute proposed here addresses that gap directly through an instrument the Governor has indicated, both through her executive order and through her signature on L.D. 713, she is prepared to support. It produces a continuous record of who is doing what compute where, on what scale, with what grid consequences, at whose expense. It preserves the option of a future moratorium on the basis of evidence rather than projection. It does not foreclose. It illuminates, and it documents, and it preserves the option to do more.