

Supporting Report: Title VI Policy Addendum

I. Overview and Motivation:

When an AI product causes harm, two distinct questions arise: who, if anyone, is legally responsible, and whether that liable party has the resources or insurance coverage to satisfy the claim. The two must operate in tandem. A defendant who is legally liable but financially unable to satisfy a judgment is, from the claimant's perspective, no different from one who is not liable at all. For most dangerous industries in the United States, such as oil tankers, nuclear reactors, and offshore drilling platforms, Congress has long paired liability rules with mandatory financial-responsibility requirements that condition market access on the demonstrated capacity to pay for foreseeable harm. [The Bipartisan AI LEAD Act, S. 2937](#), implements the first half of that framework for AI: it establishes a federal products-liability regime for AI developers and deployers, including strict liability for unreasonably dangerous covered products. It does not do the second.

Our policy addendum, Title VI, supplies what is missing. It is a four-section amendment that conditions AI market access on demonstrated ability to pay for foreseeable harm, pools catastrophic risk across frontier developers through a federally-indemnified fund, and clarifies how existing insurance policies apply to AI. By grounding these standards in a bipartisan legal framework, this addendum transforms AI safety from abstraction to a pro-worker guarantee with a clearer path to enactment.

This is critical for making the AI LEAD Act operational. A liability rule without a collection mechanism is a “paper right”; S. 2937 currently lacks financial responsibility requirements, provisions for catastrophic loss, or enforcement hooks against foreign developers. [In Garcia v. Character Technologies \(2025\)](#), wrongful-death claims proceeded against a three-year-old startup and its founders (they were recently [settled](#)). Whether the victims obtain meaningful compensation may depend substantially on corporate structure and insurance as well as the strength of their legal theories. To avoid the “collectibility problem” seen in past industrial disasters, Title VI ensures that strict liability is backed by a demonstrated capacity to pay.

Beyond compensation, Title VI also pushes developers to build safer AI in the period before courts and federal agencies finish defining what “reasonable care” requires. Consider the following causal chain: an insurer cannot issue coverage without pricing the risk, and in many cases cannot price the risk adequately without some assessment of the developer's safety practices. A developer who cannot demonstrate systematic red-teaming, pre-deployment capability evaluations, deployment safeguards, and the auditable governance set out in [NIST AI 100-1](#) and [ISO/IEC 42001:2023](#) may face materially higher premiums, narrower coverage, or an inability to obtain coverage on acceptable terms—and under Title VI, absent a qualifying form of financial responsibility, there is no lawful market access above Tier 1. Insurance becomes a

parallel regulator that does not merely compensate harm after the fact, but prevents it *ex ante* by making safety legible, auditable, and economically unavoidable.

II. The Governance Gap: What S. 2937 Does Not Solve

The AI LEAD Act has three vulnerabilities that liability doctrine cannot solve on its own. Each corresponds to a historical regulatory failure that Congress has previously addressed through financial-responsibility requirements.

First, *judgment-proofness at the developer tier*. [Section 302\(a\)](#) provides that when a developer is “not subject to the court’s jurisdiction, is insolvent, or cannot otherwise be made to answer for the harm,” the deployer may be held liable in the developer’s place. This allocates loss between defendants but does nothing for the claimant when both defendants are financially incapable of satisfying a judgment, or when damages exceed their combined assets. The AI industry is unusually exposed to this failure mode. Frontier developers are often young, venture-backed companies whose balance-sheet assets are small relative to the harms their products can foreseeably cause; Character Technologies, the lead defendant in *Garcia*, was incorporated in 2021 and [valued at roughly \\$1 billion in its August 2023 funding round](#)—which is substantial by venture standards but potentially trivial against the aggregate exposure of wrongful-death and class claims proceeding in multiple jurisdictions. Deployers who integrate third-party models face a structurally similar problem: a small business deploying a foundation model through an API may be the nearest solvent target for an injured claimant even though it exercised no control over the model’s design. Congress confronted a closely analogous collectibility problem in the oil-spill context and responded, in [section 1016](#) of the [Oil Pollution Act](#) of 1990, by requiring responsible parties to demonstrate evidence of financial responsibility in amounts sufficient to cover their potential liability as a condition of lawful operation for covered vessels and facilities. The statute transformed strict liability from a theoretical right into an enforceable one.

Second, *unenforceability against foreign developers*. Title IV of the AI LEAD Act requires foreign developers to designate an agent for service of process before making a covered product available in the United States, and authorizes the Attorney General to seek injunctive relief for noncompliance. That requirement facilitates service and enforcement, but by itself does not ensure the existence of reachable assets to satisfy a judgment. If a developer’s assets are entirely offshore, as is the case with many rising non-U.S. firms, victims are left with “empty judgments.” We should adopt the maritime approach, which requires foreign entities to post security attachable to U.S. assets or face a denial of market access.

Third, *correlated-risk accumulation at the frontier*. Section 101(a)(4) imposes strict liability on developers for unreasonably dangerous defective products, an appropriate deterrent for systems whose foreseeable risks cannot be fully anticipated at the time of release. But strict liability applied to a frontier model deployed through an API to thousands of downstream actors produces a distinctive loss-aggregation pattern: a single design defect, or a single unanticipated emergent

behavior, can generate claims across every deployment simultaneously. [Gabriel Weil has argued](#) that this is the central reason tort law alone cannot internalize AI's catastrophic externalities: the expected value of the worst-case outcomes is too large for any individual firm's balance sheet, and private insurance markets will not voluntarily provide capacity at the levels required. The bill's strict-liability rule correctly identifies who should bear the residual loss; it provides no mechanism by which that loss can actually be borne when it materializes at scale.

III. How the Proposal Works in Practice

This section describes the operational mechanics of Title VI. To summarize: § 601 requires developers to carry insurance or an equivalent financial guarantee scaled to their risk exposure; § 602 adds a pooled backstop for catastrophic events that exceed private capacity; § 603 resolves two structural problems in the insurance market itself, ambiguous coverage in existing policies and the absence of a clear compliance benchmark for underwriters; and § 604 handles technical provisions.

A. § 601: Tiered financial responsibility. Section 601 places each developer into one of four tiers based on two alternative triggers: the training compute used to build the model and the annual revenue the model generates. A developer falls into a higher tier if *either* threshold is exceeded. This disjunctive trigger is essential because compute and revenue measure different kinds of risk. Compute is a proxy for capability and therefore for the severity of potential harm; revenue is a proxy for deployment scale and therefore for the breadth of potential exposure. (We discuss potential drawbacks to this approach later on.) A heavily-deployed but low-compute chatbot, such as the systems at issue in [Garcia](#) and [Raine](#), would trigger a financial-responsibility obligation under the revenue prong even though it would fall below every compute threshold in current or proposed AI legislation. Conversely, a frontier model with limited commercial deployment would trigger the compute prong.

Section 601 permits four forms of financial responsibility: commercial liability insurance, a surety bond, an irrevocable letter of credit, or qualification as a self-insurer upon demonstration of tangible net worth of at least ten times the required minimum. This is adapted in part from the Oil Pollution Act and its implementing [financial-responsibility regime](#), which has administered evidence-of-financial-responsibility requirements for vessel operators since the early 1990s. The logic is that different firms at different stages of development have different optimal capital structures. A venture-stage startup may find commercial insurance the only viable option; a mature developer may prefer a letter of credit or self-insurance; large incumbents such as Google or Microsoft could likely qualify as self-insurers under a ten-times-net-worth test. Forcing all firms into a single instrument would either price smaller developers out of the market or waste capital at larger ones. The § 601(c) open-source exemption excludes developers who release model weights and architecture under a license permitting unrestricted use and redistribution, except where a Tier 4 systemic designation applies. Section 601(a) requires foreign developers to maintain financial responsibility in a form enforceable against assets located in the United States,

and authorizes the Attorney General to direct compute providers, payment processors, and content delivery networks to cease facilitating availability of noncompliant covered products, an intermediary-directed enforcement model that is analogous, though not identical, to [OPA 90's clearance-denial mechanism](#) and in [Treasury OFAC practice](#).

B. § 602: Advanced AI Compensation Fund. Section 602 establishes a layered compensation mechanism for qualifying AI incidents, events that cause harm to more than 500 persons, or aggregate damages exceeding \$500 million, as certified by the Attorney General. This structure closely parallels the two-layer financing system that the [Price-Anderson Act of 1957](#) created for nuclear liability, which has successfully underwritten the commercial nuclear power industry for nearly seven decades [without interruption](#). First, Tier 4 developers must carry \$1 billion in private insurance or self-insurance. If a disaster exceeds that amount, the second layer kicks in: a “pay-as-you-go” shared pool where all top-tier developers contribute up to \$500 million each. This [“retrospective” design](#) is vital because it provides massive emergency funding without forcing companies to tie up capital for risks that may never happen. Third, if both layers are exhausted, the U.S. Treasury provides a backstop of up to \$30 billion to ensure victims are compensated.

There are two specific design choices which we will expand on here for clarity. First, the Fund’s “secondary character,” specified in § 602(b), means it pays only to the extent a claimant’s recovery is not otherwise satisfied under Title I. This design preserves the primary tort system as the main deterrent while using the Fund as a compensation backstop rather than a liability shield. Second, the qualifying-incident threshold (500 persons or \$500 million) is set lower than the catastrophic-risk threshold in [California’s SB 53 \(more than 50 deaths/serious injuries or \\$1 billion\)](#). The lower threshold is deliberate: SB 53 concerns [transparency disclosures](#) to the California Office of Emergency Services, where over-triggering creates administrative burden, whereas the Fund concerns claimant compensation, where under-triggering leaves victims without recourse. The \$30 billion federal indemnification ceiling is substantial but finite. It is not intended to cover truly catastrophic or existential risks, which [Professor Gabriel Weil has argued](#) because compensation becomes meaningless at scales approaching societal collapse.

C. § 603(a): Silent-AI affirmation. Section 603(a) addresses a problem that is forming in the commercial insurance market independent of the AI LEAD Act. Most existing commercial general liability, errors and omissions, professional liability, directors and officers, and cyber policies were drafted before generative AI became widespread and neither expressly cover nor expressly exclude AI-related harm. This is the same “silent” coverage problem that produced the cyber insurance crisis of 2017–2020, culminating in the [\\$100+ million coverage dispute between Mondelez International and Zurich over NotPetya losses](#) and prompting [Lloyd’s of London to mandate, in Market Bulletin Y5258 \(July 2019\) and its January 2020 implementation order, that all syndicate policies either expressly affirm or expressly exclude cyber coverage](#). Section 603(a) applies the same logic to AI: insurers have 24 months to state, in writing, whether their policies

cover AI-related harm. Importantly, under this proposal, silent policies are construed in favor of coverage. A default-exclusion rule would protect insurer balance sheets at the expense of policyholders who reasonably believed they were covered; a default-coverage rule forces active drafting, which is the point of the 24-month compliance window.

D. § 603(b): Conditional safe harbor. Section 603(b) creates a rebuttable presumption that a developer exercised reasonable care if it substantially complied, at the time the product left its control, with the [NIST AI Risk Management Framework \(NIST AI 100-1\)](#) and its [Generative AI Profile \(NIST AI 600-1\)](#), or with [ISO/IEC 42001:2023](#). This serves two purposes simultaneously. First, it gives underwriters the auditable technical benchmark they need to price negligence-based coverage. Insurance markets cannot price risks without standards, and the AI industry currently has no widely-adopted equivalent to the building codes and safety standards that underpin property and casualty insurance in other sectors. Endorsing NIST RMF and ISO 42001, both the product of extensive multi-stakeholder processes, provides such a standard without forcing Congress to specify technical requirements that will be obsolete within years. Second, and more importantly, the presumption is carefully limited. It applies only to the negligence theories in § 101(a)(1)–(2), not to the express-warranty theory in § 101(a)(3) or the strict-liability theory in § 101(a)(4). Strict liability under § 101(a)(4) is justified precisely on the theory that developers should bear residual loss *even when care was exercised*; allowing framework compliance to rebut strict liability would meaningfully hamper the bill’s deterrent function for catastrophic-risk-generating development.

E. § 604: Conforming provisions. Section 604 performs three ‘housekeeping functions’. It preserves claimants’ rights under Title I, confirming that compliance with Title VI does not limit any claimant’s right of recovery. It indexes all dollar figures to the Consumer Price Index for All Urban Consumers, using [the same CPI-U adjustment mechanism that 33 U.S.C. § 2704\(d\)\(4\) applies](#) to OPA 90 liability limits, without which the \$10 million, \$100 million, and \$1 billion thresholds would erode in real terms over the life of the statute. And it redesignates the original Title V (Effective Date) as Title VII and the original § 501 as § 701, just to keep the bill’s existing effective-date language operative.

IV. Technical Foundations and Assumptions

Title VI rests on four technical assumptions, which we will unpack here.

The first assumption is that training compute is a usable proxy for model capability and therefore for potential harm. Compute thresholds have been adopted by several prominent AI regulatory instruments that have attempted to calibrate obligations by model capability: [EO 14110 \(10²⁶ FLOPs\)](#), [SB 53 \(10²⁶\)](#), [EU AI Act Article 51 \(10²⁵\)](#). But the proxy is increasingly suspect. [Fenwick’s 2024 analysis](#) documents how mixture-of-agents architectures and distributed training can produce frontier-level capabilities while keeping any single training run below the threshold, and DeepSeek’s 2025 models have [demonstrated](#) that compute-cheap models can nonetheless

exhibit high-risk behaviors including susceptibility to jailbreaking. Our response is the disjunctive compute-or-revenue trigger: a model that evades compute thresholds through architectural innovation still triggers financial responsibility obligations if it achieves commercial scale.

The second assumption is that AI risk is actuarially tractable at middle tiers but requires pooling at the frontier. This is the proposal's deliberate position in the current scholarly debate. [Anat Lior's empirical work](#) documents an AI liability insurance market that is nascent but increasing in scope, with specialty insurers like [Testudo](#) launching in 2025 and major carriers developing bespoke products. [Daniel Schwarcz and Josephine Wolff have argued](#) that cyber insurance failed to produce the security outcomes once hoped for, because insurers priced legal risk rather than technical risk and funded defense lawyers rather than security investments. Tiers 2 and 3 rely on private market discipline, where AI risks are specific enough to underwrite; Tier 4 requires participation in a compulsory pool, because correlated risk at frontier scale defeats private insurance capacity in the same way it defeated private nuclear insurance before 1957. There is [credible evidence](#) that the cyber-insurance failure was in part a failure of underlying liability clarity, which Title I of the AI LEAD Act now supplies.

The third assumption is that substantial compliance with the NIST RMF is verifiable enough to serve as a legal presumption. From our analysis, NIST's framework is outcome-oriented and non-prescriptive rather than a formal checklist, which complicates adjudication but also preserves flexibility. This is a feature for risk management, but it does complicate adjudication. Title VI accepts the tradeoff for two reasons. First, the structure is sufficiently specific that substantial compliance can be documented and audited, even if the substantive adequacy of implementation remains a factual question. Second, the presumption is rebuttable, not conclusive: a plaintiff can show that implementation was inadequate *in the specific circumstances giving rise to the harm*. This preserves courts' fact-finding role and reduces the risk that formal compliance alone becomes a shield against otherwise meritorious claims.

The fourth and most consequential assumption is that strict liability must remain outside the safe harbor. The [Restatement \(Third\) of Torts: Products Liability § 2\(a\)](#) treats manufacturing defects as generating strict liability on the theory that the manufacturer is the least-cost avoider of residual loss even when all reasonable care was exercised. Professor [Gabriel Weil has extended this logic to catastrophic AI risk](#): the externalities of frontier AI development are sufficiently severe and sufficiently novel that fault-based deterrence is insufficient, and strict liability is the only doctrinal tool capable of forcing developers to internalize the full expected cost of the systems they release. The carve-out is what lets Title VI extend insurance infrastructure into fault-based liability without disturbing the bill's fault-less core.

V. Why the Chosen Mechanisms and Thresholds Are Appropriate

Here we expound on three design choices in Title VI: the specific threshold values, the specific dollar amounts, and the open-source carve-out. Each was chosen against a field of plausible alternatives.

The thresholds are borrowed from existing legislation for harmonization. The 10^{26} -FLOP ceiling tracks the compute threshold used in [EO 14110](#) and [California's SB 53](#); the 10^{25} midpoint tracks [Article 51 of the EU AI Act](#); the 10^{23} floor captures the next order of magnitude down, closing a coverage gap that the [California Report on Frontier AI Policy \(June 2025\)](#) explicitly identified—namely, that models below SB 53's ceiling may still pose significant catastrophic risk. Borrowing thresholds lowers compliance novelty, reduces international regulatory friction, and positions the regime to evolve in step with the reassessment schedules already built into those instruments (SB 53's first reassessment is [scheduled for 2027](#)). The \$500 million large-developer revenue threshold in SB 53 informs the scaling logic for Title VI's revenue prongs, which step down proportionally to \$100 million, \$10 million, and zero.

The dollar amounts are proportionate to industry capacity, and not excessively punitive. Tier 2's \$10 million tracks standard commercial general liability program limits. Tier 3's \$100 million sits below the [\\$500 million primary layer that American Nuclear Insurers has provided since 2024](#) for the nuclear industry. Tier 4's \$1 billion is well below the overall Price-Anderson compensation pool, which now [exceeds \\$16 billion](#). For context, [U.S. private AI investment reached \\$109.1 billion in 2024](#) and [\\$285.9 billion in 2025](#) per the Stanford AI Index; the financial responsibility thresholds represent a small fraction of the capital already deployed in the sector. This is important: financial responsibility requirements that exceed what firms can realistically capitalize function as de facto prohibitions, and Title VI is not designed to prohibit AI development. It is designed to ensure that developers whose systems cause harm are capable of paying for it.

The open-source carve-out in § 601(c) is one of the proposal's most politically significant design choices. [SB 1047 was vetoed in September 2024](#) and its liability and control provisions were [widely criticized](#) as discouraging open-weight or open-source release. SB 53's [narrower approach reflected those lessons](#). Section 601(c) threads the needle: developers who release their model weights and architecture under a permissive license, without monetizing the release beyond reasonable cost recovery, are exempt from financial responsibility obligations entirely, *except* where the model independently crosses the Tier 4 systemic-risk threshold. This preserves the research, competition, and safety benefits of open models (which [Gabriel Weil has argued](#) include downstream alignment research that cannot be performed behind APIs) while preventing the carve-out from becoming a loophole for frontier-scale systemic risk. The non-attribution rule in § 601(c), third-party deployer revenue is not imputed back to the releasing developer, is essential to this design: without it, a developer releasing weights would face open-ended liability for downstream monetization they neither controlled nor anticipated.

Two smaller choices are worth noting. The silent-AI default-coverage rule in § 603(a) runs pro-claimant rather than pro-insurer because a default-exclusion rule would protect insurer balance sheets at the expense of policyholders who reasonably believed they were covered, the outcome [Airmic warned against during the Lloyd's rollout](#). The Fund's secondary character in § 602(b), mirroring [42 U.S.C. § 2210\(e\)](#)'s layered structure, preserves the primary tort system as the principal deterrent. So the Fund compensates, but does not substitute for liability.

VI. Interpretation, Implementation, and Enforcement

There are three relevant implementation questions:

First, *administering agency*. The Attorney General administers Title VI because [§ 401 of the bill](#) already implicitly routes foreign-developer registration through DOJ. Administration of filings and the registry under § 601 parallels the [National Pollution Funds Center's administration of OPA 90 Certificates of Financial Responsibility](#), which is housed within a single operational unit rather than distributed across a new agency. Title VI does not require new agency creation.

Second, an *enforcement cascade*. Noncompliance triggers a tiered response. First, civil penalties under [§ 301\(2\)](#), which already authorizes the Attorney General to seek them. Second, injunctive relief under the same provision, prohibiting the developer from making the product available in interstate commerce pending cure. Third, for foreign developers specifically, § 601(a) authorizes cease-facilitation orders directed to compute providers, payment processors, and content delivery networks. This is the most structurally novel enforcement mechanism in Title VI, but it is modeled on two well-established precedents: [OPA 90's clearance-denial mechanism under 33 CFR § 138.170](#), which denies port access to noncompliant foreign vessels, and [Treasury OFAC's practice of directing financial intermediaries](#) to freeze or decline transactions with sanctioned parties. The common principle is that jurisdiction over intermediaries operating on U.S. soil supplies enforcement leverage when direct jurisdiction over the primary actor is unavailable.

Third, *effective date*. An 18-month effective date, deferred from enactment, aligns with the [24-month silent-AI window](#) and with SB 53's January 1, 2026 implementation date. This sequencing was deliberate: insurers have six months after § 603(a) takes effect to finalize coverage language before § 601's financial responsibility requirements bind developers.

VII. Limitations

First, compute is a leaky proxy for capability and the leaks are widening. [Fenwick's 2024 analysis](#) documents how mixture-of-agents architectures can achieve frontier capabilities by stitching together sub-threshold models, and DeepSeek's 2025 releases [demonstrated](#) that algorithmic efficiency can produce high-risk behaviors at compute levels well below 10^{26} FLOPs. Title VI's disjunctive revenue trigger partially mitigates this, as a heavily deployed but relatively low-compute chatbot, such as the systems at issue in Garcia and Raine, would likely trigger a financial-responsibility obligation under the revenue prong even though it would fall

below the principal compute thresholds in current or proposed AI legislation. The proposal does not include a delegated reassessment authority; adding one would be a sensible amendment but raises its own constitutional questions about nondelegation.

Second, the \$30 billion federal indemnification ceiling is deliberately bounded. Professor [Gabriel Weil has argued](#) that truly catastrophic AI risk, at the scale that could approach civilizational disruption, cannot be internalized through compensation mechanisms because compensation becomes meaningless at sufficient scale, and requires a punitive-damages supplement designed to pull forward expected liability into cases of practically compensable harm. Title VI does not attempt to address existential-scale risk. It addresses the substantial band of foreseeable AI harm that lies between individual tort claims and civilizational collapse, and leaves the tail problem for distinct policy instruments.

VIII. Risks and Mitigations

The principal risks are moral hazard and capacity. Moral hazard describes the tendency of an insured party, once shielded from the full consequences of a harmful outcome, to take fewer precautions against causing it. In this context, the concern is that a developer whose catastrophic losses are backstopped by the Fund will invest less in safety than a developer fully exposed to tort liability. This critique has followed Price-Anderson [since its passage](#). Three features mitigate it. First, the Fund is secondary to primary recovery under Title I, meaning developers remain the first line of financial responsibility. Second, strict liability under § 101(a)(4) is preserved outside the safe harbor, meaning compliance with technical standards does not discharge liability for unreasonably dangerous products. Third, § 604 explicitly disclaims any limit on claimant recovery, meaning the Fund operates as a compensation backstop rather than a liability shield.

The capacity concern is more specific. The entire U.S. cyber insurance market, the most analogous commercial line, was approximately [\\$9.14 billion in direct written premium in 2024](#), with the market [recording its first-ever decline](#) that year. Tier 4's \$1 billion minimum, applied across even a handful of frontier developers, would approach the limits of current standalone commercial capacity. This is precisely why § 602 uses a retrospective premium model rather than requiring *ex ante* capitalization: capacity is called only when needed, spread across the full population of Tier 4 participants, and reinforced by federal indemnification above the private layer.

IX. Conclusion

Title VI does not replace S. 2937's liability architecture. It makes that architecture collectible, enforceable against foreign actors, and actuarially tractable. The amendment draws on three of the best-evidenced statutory precedents for governing heavy-tail risk: the [Price-Anderson Act's layered compensation structure](#), the [Oil Pollution Act's financial responsibility regime](#), and [Lloyd's silent-cyber mandate](#). Each is adapted to the technical and constitutional realities of AI rather than imported wholesale.

The result answers each of the structural audiences an AI liability regime must satisfy. Claimants obtain recovery when their judgments would otherwise be uncollectible. Developers obtain certainty about what capitalization is expected of them and what standards, when substantially met, establish a rebuttable presumption of reasonable care. Insurers obtain the technical benchmark they need to price coverage and the coverage clarity they need to avoid silent-AI accumulation losses. And the public obtains the catastrophic-risk backstop that private markets cannot provide alone. S. 2937 establishes who is liable for AI-related harm. Title VI ensures that liability can be paid.

More fundamentally, Title VI solves the central institutional defect of AI tort law: liability without payment is not deterrence, not compensation, and not governance. By tying ex ante capitalization, insurance-market clarity, and a public excess layer to the underlying liability scheme, Title VI converts formal liability into a functioning compensation system. In high-consequence domains, law must do more than announce responsibility after disaster; it must ensure, before disaster, that responsibility can be borne. Title VI supplies that missing institutional bridge. Without it, S. 2937 risks becoming a declaration of principle. With it, the bill becomes a workable regime for allocating and absorbing AI catastrophe risk.

Appendix

Since the academic literature on this subject is not fully comprehensive and AI risk changes so rapidly with every new model update, we spoke to experts in the field on their perspectives on the future of AI insurance. We conducted interviews with Professor Anat Lior at Drexel University, focusing on her paper in the Harvard Journal of Law & Technology on the role of insurance in AI regulation and Alexander Choy, who is an insurance underwriter at Relm (insurance firm focused on AI risk), and Christian Trout, who's a member of the policy team at AIUC (a firm that seeks to certify and insure AI agents). We are including the transcripts of these interviews below, since they informed many of our policy prescriptions in the bill.

1. Interview Notes with Professor Anat Lior, Assistant Professor of Law at Drexel University

Q: Your analysis of the “green lumber problem” was really interesting. Curious to know what your perspective is on how to truly evaluate litigation risk for these models? Is it a combination of technical analysis on things like hallucination, chain of thought reasoning to create a quantified risk profile or some other form of auditing?

A: This analysis steamed with conversations I had with Testudo who are opposed to just relying on audits and evaluations. They are using current AI litigation cases to help them predict better liability outcomes. Not sure that this is a perfect parameter, but it should be traditionally better than mere audits. I agree that a deeper understanding via the underwriting process is needed, but currently there is a heavy reliance on existing litigation as a means to predict future litigation.

Q: On the same note, what benchmarking metrics did your research participants identify as being most closely correlated with actual loss data?

A: They did not name specific benchmark metrics unfortunately. I can put you in touch with Armilla AI if you want because I think they will be happy to talk about this.

Q: Your research focused on 'silent AI' as a major source of uncertainty for traditional policy providers. Are there any specific industries that are more worried about this unintended exposure and what steps are companies taking to de-risk in the eyes of insurance providers (if any)?

A: I think medical malpractice is a unique industry that is very worried about this. They are working with AI exports and being very cautious regarding what AI based products and services they are implementing into their workflow. Other industries relate to IP issues and they are mostly increasing awareness of employees by crafting specific AI policies for companies.

Q: What would push insurers to stop relying on this silent coverage and start demanding benchmarks from these companies?

A: More claims. I think that they don't have a clear incentive yet to do anything about this until they start losing more money via claim reimbursement.

Q: There's a couple companies in the market that are trying to tackle this issue (ex: Fulcrum, Relm, Armilla, etc) - any thoughts on their specific approach of either 1) helping insurance brokerage companies automate manual work or 2) independent insurance brokerages writing policies with AI risk in mind.

A: I haven't been thinking enough about this aspect of AI and insurance. I focus in my work on how liability policies will cover AI, and less on AI will make the insurance industry more efficient. Happy to talk more about this and put you in touch with people who have done more work on this.

2. Interview Notes with Alexander Choy, Underwrites at Relm AI

Q: Who is currently dealing with AI insurance pricing in the marketplace?

A: Relm is the only one in the market place that deals with AI insurance pricing. There are other companies that are attempting to provide AI insurance but most insurance right now is just basic cyber insurance: If there is no human in the loop - will just have bad coverage and tons of exclusions - would probably just have coverage for stealing IP/ransomware against the models.

Q: What specific questions do they ask brokers before insuring AI companies?

A: [He provided the exact questionnaire that they send companies before insuring them]

- Please provide details of your AI model training methods including sources of training data.
- What specific risk management practices and technical safeguards has the company implemented to prevent its AI solution from inadvertently or unlawfully disclosing sensitive or protected information?
- Has the Applicant implemented a formal dataset governance process to verify that all training data is lawfully obtained and properly licensed?
- Has the Applicant implemented documentation controls for all datasets used in model training, including audit trails of data sources?
- Has the Applicant implemented licensing agreements with third-party rights holders (e.g., stock image/music libraries, text archives), and are these agreements regularly reviewed for compliance?
- Has the Applicant implemented controls to exclude pirated, paywalled, or otherwise unauthorized content (e.g., subscription media, copyrighted books) from its training datasets?
- Has the Applicant implemented review procedures for open-source datasets to confirm that their terms of use allow for commercial AI training?
- Has the Applicant implemented processes to periodically revalidate training datasets to ensure ongoing compliance with evolving copyright and privacy regulations (e.g., EU AI Act, U.S. state privacy laws)?
- Has the Applicant implemented third-party audits or legal reviews of its dataset sourcing and training model practices?

Q: How does the insurance industry approach AI given it is typically backward-looking?

A: Insurance is very backward looking that is mostly empirical and in order to do AI (no data point that you can look at and you have to be more proactive). Indicative of the insurance market, they have to price based on the market (other insurance companies). Also, since they don't have the technical talent to deeply evaluate the model architecture, most of the questions they ask are very surface level.

Q: What is the current state of losses and policy pricing?

A: So far not seeing losses there's no history but no one wants to buy an expensive policy — right now it's literally looking at revenue and cyber evals (soft market for the insurance). Need something catastrophic to happen to go from soft market to hard market for the insurance market to reset. Otherwise, the policies will still be “too cheap” and not actually reasonable.

Q: Are AI models fully insured under current standards? This question is inspired by the recent ElevenLabs announcement that their partnership with AIUC has led to a “fully insured” voice agent.

A: No AI model is fully insured. If they are claiming this, that means all liability is put on the individual and that's not how insurance litigation works. People that are smart enough to evaluate models are not looking at models - interested in AI and understanding at the foundational level.

Q: How are emerging risks like hallucinations being handled?

A: He did not know what hallucination was and said that they don't cover it - huge downstream risk. Most of the AI companies are very uninsured... There are privacy acts that try to solve some of it but it isn't enough.

3. Interview Notes with Member of Policy Team at AIUC (AI Underwriting Company)

Q: How does liability frameworks today implicate foundation model developers?

A: Liability framework proposes strict (not exclusive) liability for foundation model developers on catastrophic losses >\$1B, using "least cost avoider" principle; allows pursuing downstream developers alongside foundation models (e.g., if Anthropic-powered Goldman agent causes \$2-4B damage). The Insurance industry is positioned to regulate AI safety similar to its historic role in auto (airbags, crashworthiness) and electrical appliances (UL standards), creating market enforcement and trust mechanisms that accelerate adoption.

Q: How does AIUC adjust premiums compared to traditional insurance brokers?

A: Red teaming + AIUC1 standard certification adjust premiums ~30%, with causal risk modeling (FMEA-style failure analysis borrowed from nuclear/aerospace) replacing traditional actuarial approaches; current customers include Intercom, ElevenLabs (application layer only, not OpenAI/Anthropic).

Actuarial models use simple historical approaches, applying past accident rates to future products without examining specifications or red teaming. Red teaming serves as middle ground between pure actuarial models and full causal modeling through stress testing.

Q: There is a massive coverage gap with frontier models/companies (OA/Anthropic) having only ~\$200M coverage vs. multi-billion exposures. Both have faced multi-billion dollar lawsuits without coverage - what's the plan?

A: They need billions in coverage, far beyond current market capacity. Catastrophic cyber risk is a cautionary example: only approximately 20% of NotPetya losses were covered despite decades of cyber attacks. Proposed solution: frontier AI companies form mutual insurance co-op (co-operative, not-for-profit owned by policyholders) rather than relying on government mandate or traditional insurers filling the gap voluntarily.

AIUC is an Managing General Agent sitting between policyholders and insurers as specialized underwriters. MGAs help price coverage without holding capital; insurers retain capital and risk-bearing function.